Conversiono

AI-Powered Behavioral Intelligence for E-Commerce

Investor & Vision Document

March 2026 | Confidential

Executive Summary

The opportunity in one page

Convorsino is building the world's first conversion psychology operating system for e-commerce. While existing analytics tools tell merchants what users did — clicks, bounce rates, page views — they are completely blind to why users hesitated, what psychological friction blocked a purchase, and how a buying decision was actually formed.

We solve this by combining a lightweight JavaScript SDK with a proprietary psychology-driven intelligence layer powered by machine learning. The result is a platform that detects the hidden emotional friction killing conversion, explains it in plain language grounded in behavioral science, and recommends specific, evidence-based interventions — all in real time.

€45B+ Global e-commerce analytics market (2026)	**2.86%** Average global e-commerce conversion rate — the rest is lost revenue
70% Of online shopping carts are abandoned before purchase	**€18B+** Revenue lost annually to preventable checkout friction (EU alone)

What makes Conversiono different

Every competitor — Hotjar, Contentsquare, FullStory, Optimizely — operates at the behavioral surface. They show you heatmaps of where users clicked. They show you session replays. They show you which A/B test variant won. None of them can tell you that a user experienced price anxiety at the shipping cost field driven by loss aversion, that this specific psychological state accounts for 34% of checkout abandonment on your shop, or that the highest-probability fix — based on your own experiment history and 23 comparable experiments in our database — is reframing the shipping cost as a savings threshold.

That depth of insight is what Conversiono delivers. And it is only possible because of the four-layer intelligence architecture at the heart of our product — a system that connects real-time behavioral signals, a proprietary psychology framework built on decades of academic research, an explainable machine learning ensemble, and a compounding experiment memory that grows more valuable with every merchant and every test.

The ask

Raising	€350,000 — Pre-Seed
Use of funds	Product development (50%), team (30%), go-to-market (20%)
Target close	Q2 2026
Equity offered	12–15% (negotiable)
Target milestone	10 pilot merchants, working ML model, €5k MRR by Month 6

The Problem

Why e-commerce analytics is fundamentally broken

The surface-level trap

The e-commerce analytics industry has spent two decades optimizing for the wrong thing. It has become extraordinarily good at measuring what users do and almost entirely incapable of explaining why they do it. This is not a minor gap. It is a fundamental structural failure that costs the global e-commerce industry tens of billions of euros every year.

Consider what happens when a shop owner notices a 68% abandonment rate on their checkout page. They open Google Analytics and see that users are leaving at step 3. They open Hotjar and see a heatmap showing most clicks are concentrated near the price field. They run three A/B tests over the next two months and find that variant B performs slightly better. They still have no idea why users are abandoning, which means they have no systematic way to fix it. They are optimizing in the dark.

> *"We had a 71% checkout abandonment rate. We tried 11 A/B tests over 6 months. Some worked a little, some made it worse. We never understood what was actually going wrong." — Mid-size Shopify merchant, €3.2M annual revenue*

What existing tools see vs what is actually happening

What analytics tools see	What is actually happening psychologically
User bounced at checkout step 3	User experienced loss aversion when shipping cost appeared unexpectedly
User spent 45 seconds on product page	User was in price comparison mode, anchoring against alternatives
User clicked Add to Cart then removed item	User experienced approach-avoidance conflict — desired item, feared regret
User rage-clicked the CTA button	User expected a modal to open — broken expectation created frustration
High scroll depth, low conversion	Decision fatigue — too many choices overwhelmed the decision process

Why this problem is growing, not shrinking

Three converging trends are making behavioral blindness more expensive every year. First, customer acquisition costs in e-commerce have risen by over 60% in the past three years as paid social advertising becomes increasingly competitive — meaning every lost conversion is more expensive to replace. Second, the explosion of mobile commerce has created entirely new behavioral patterns that desktop-era analytics tools were not designed to measure. Third, consumer psychology is becoming more sophisticated: modern shoppers are more aware of persuasion techniques and more resistant to them, making psychological authenticity — the ability to address real friction rather than apply surface-level tricks — increasingly important.

The shops that will win the next decade are those that genuinely understand their customers' psychology. The tools to do that do not yet exist at scale. Conversiono is building them.

The Solution

A conversion psychology operating system

The product in one sentence

BehaviorIQ is an AI-powered behavioral intelligence layer that sits on top of any e-commerce website, detects the psychological friction blocking conversion in real time, explains it in plain language grounded in behavioral science, and recommends evidence-based interventions — all without requiring any technical expertise from the merchant.

Three connected capabilities

1. Real-time behavioral intelligence

A lightweight JavaScript SDK (under 15KB, loads asynchronously) installs on any website in under five minutes. It captures a continuous stream of behavioral signals — mouse velocity, hesitation patterns, scroll behavior, click dynamics, exit intent — and processes them through a four-layer intelligence architecture that translates raw browser events into named psychological states: price anxiety, decision fatigue, trust gap, loss aversion, approach-avoidance conflict.

The merchant dashboard shows not just where users are abandoning, but why — with friction scores per page element, psychological state labels, and session-level abandonment risk scores updating in real time.

2. Experiment memory with semantic search

Every A/B test a merchant runs — whether through BehaviorIQ, Google Optimize, VWO, Optimizely, or a spreadsheet — can be stored in a structured knowledge base. Results are automatically tagged with the psychological mechanism involved, the friction type addressed, and the measured conversion impact. Over time, merchants build an institutional memory of what works for their specific audience.

The knowledge base is searchable via natural language semantic search. A merchant can ask 'what have we tested on checkout price presentation?' and receive a ranked list of relevant past experiments with outcomes, psychological reasoning, and predicted applicability to their current friction profile.

3. Friction-first test recommendations

Unlike all existing A/B testing platforms, which require merchants to come up with their own hypotheses, BehaviorIQ generates test recommendations from the bottom up. Because the intelligence layer already knows exactly where friction is occurring and what psychological mechanism is driving it, the suggestion engine can produce specific, justified, prioritized recommendations — ranked by predicted conversion impact and grounded in both academic research and the merchant's own experiment history.

What competitors offer vs what we offer

Capability	Hotjar	Contentsquare	Optimizely	BehaviorIQ
Behavioral heatmaps	Yes	Yes	No	Yes
Session recordings	Yes	Yes	No	Yes
A/B testing	No	Limited	Yes	Via integrations
Psychological friction detection	No	No	No	Yes
Real-time abandonment prediction	No	No	No	Yes
Explainable ML (SHAP)	No	No	No	Yes
Experiment memory & search	No	No	Limited	Yes
Friction-first test suggestions	No	No	No	Yes
Psychology-grounded explanations	No	No	No	Yes

The Intelligence Layer

The technical heart of BehaviorIQ — a deep dive

The intelligence layer is what separates BehaviorIQ from every analytics tool in the market. It is not a single model or algorithm — it is a four-layer system in which each layer feeds the next, converting raw browser events into actionable psychological insight in under 300 milliseconds. This section provides a complete explanation of how each layer works, why it is designed the way it is, and how the layers interact.

Layer 1 — The Psychology Framework

The psychology framework is the intellectual foundation of the entire product. It is a structured knowledge base — not code, not an ML model, but a carefully researched mapping between observable behavioral patterns and named psychological states, grounded in decades of peer-reviewed behavioral economics and cognitive psychology research.

The framework answers a question that no algorithm can answer on its own: given that a user's mouse slowed to near-zero velocity near a price field for 4.2 seconds before their cursor drifted toward the top of the browser — what does that mean about what was happening in their mind?

The answer comes from three bodies of research. Daniel Kahneman and Amos Tversky's Prospect Theory establishes that humans experience losses approximately twice as painfully as equivalent gains — the behavioral signature of this is hesitation and retreat near cost-related elements. Barry Schwartz's Paradox of Choice research demonstrates that excessive options produce decision paralysis — measurable as long session duration with low progression and repeated scroll reversals. Robert Cialdini's research on social proof, scarcity, and authority maps directly to specific behavioral patterns around trust signals and urgency indicators.

Behavioral pattern observed	Psychological state	Research basis	Intervention direction
Mouse slows near price, scroll retreat	Loss aversion / price anxiety	Kahneman Prospect Theory	Reframe cost as gain or anchor against higher reference
Repeated scroll up-down near checkout	Decision uncertainty	Schwartz Paradox of Choice	Reduce options, add social proof, simplify decision
Long dwell, low progression	Decision fatigue	Baumeister ego depletion	Highlight single recommended option, reduce cognitive load
Rage clicks on non-responsive element	UI frustration	Nielsen Norman broken expectation	Fix interaction design, clarify affordance

Cursor drifts to browser top edge	Abandonment intent	Eye-tracking exit intent research	Trigger retention intervention immediately
Add to cart → remove → re-add	Approach-avoidance conflict	Lewin field theory	Reduce perceived risk — guarantee, reviews, returns policy
Fast scroll past trust signals	Trust gap	Fogg persuasion technology	Surface trust signals earlier, increase prominence

This framework is stored as a versioned configuration and is updated as new research emerges or as our own data reveals new patterns. It is the proprietary intellectual property that no competitor can replicate quickly — not because the research is secret, but because synthesizing it into a precise, production-ready behavioral taxonomy takes years of domain expertise.

Layer 2 — The Feature Engineering Engine

The feature engineering engine is the bridge between raw browser data and the numerical inputs the ML models require. Its job is to transform a stream of thousands of mouse events, clicks, and scroll actions into a compact, meaningful vector of numbers — each number capturing one specific psychological signal.

This transformation is not trivial. A mouse movement at coordinates (340, 210) at timestamp 1741203842 is meaningless in isolation. But when you observe that the mouse velocity dropped below 15 pixels per second near the element with id='shipping-cost-line' and remained there for 4.2 seconds, that is the hesitation score for that session on that element — a number with direct psychological meaning.

The engine computes features continuously, updating the session's feature vector every time a new batch of events arrives from the SDK. It maintains a stateful session object in memory so that running computations (like average velocity or cumulative hesitation time) do not need to be recalculated from scratch on every update.

Feature name	What it measures	How it is computed	Psychological signal
hesitation_score	Mouse pause duration near price/cost elements	Avg(pause_duration) × pause_frequency where velocity < 15px/s near cost elements	Price anxiety, loss aversion
scroll_retreat_count	Number of scroll direction reversals	Count(events where scroll_direction ≠ previous_direction)	Decision uncertainty, confusion
rage_click_score	Rapid repeated clicks on same element	click_count / time_on_element where clicks > 3 in 2 seconds	UI frustration, broken expectation

exit_intent_count	Cursor movements toward browser top edge	Count(events where y < 20px and velocity > 50px/s upward)	Abandonment intent
checkout_hesitation_s	Inaction time on checkout pages	Duration on checkout pages with no scroll/click activity	Price shock, commitment anxiety
price_dwell_time_s	Time cursor spent near price fields	Sum(time intervals where cursor within 40px of price element)	Price evaluation intensity
velocity_variance	Erratic vs smooth mouse movement	Standard deviation of velocity across session	Cognitive overload, anxiety level
session_duration_s	Total session length	last_event.timestamp − first_event.timestamp	Engagement depth vs fatigue

Layer 3 — The Machine Learning Ensemble

The ML layer consists of three specialized models running in parallel on each session's feature vector. Rather than attempting to solve all three prediction problems with a single model — which would reduce both accuracy and explainability — each model is trained for exactly one task, and their outputs are combined into a unified session intelligence object.

All three models are built on XGBoost, a gradient boosted decision tree algorithm chosen for three specific reasons: it performs exceptionally well on tabular data at the scale we will have in V1 (1,000–50,000 sessions), it trains in minutes on standard hardware without requiring GPU infrastructure, and it is natively compatible with SHAP (SHapley Additive exPlanations), the explainability technique that allows us to attribute each prediction to specific input features.

Model 1 — Intent Classifier

The intent classifier answers the question: what mode is this user in right now? It performs multi-class classification across four states — browsing (user is exploring with no clear intent), comparing (user is evaluating options, likely has alternatives open), deciding (user is close to a purchase decision but has not committed), and exiting (user is disengaging and likely to abandon).

This classification matters because the appropriate intervention is entirely different for each state. A user in comparing mode may respond positively to social proof and differentiation messaging. A user in deciding mode may need a friction reduction — a clearer guarantee or simplified checkout. A user in exiting mode needs an immediate retention trigger.

Model 2 — Friction Detector with SHAP Explainability

The friction detector is the most technically sophisticated component of the intelligence layer. It predicts not just that friction exists, but where it is located and what psychological mechanism is driving it — and it does so in a way that is fully explainable to non-technical merchants.

The key technical innovation here is the use of SHAP values. After the model makes its prediction, a SHAP TreeExplainer computes the marginal contribution of each feature to that specific prediction. If the friction score for a session is 0.84, SHAP tells us that checkout_hesitation_s contributed +0.31, exit_intent_count contributed +0.24, and price_dwell_time_s contributed +0.19 to reaching that score. These feature contributions are then cross-referenced with the psychology framework to produce a human-readable diagnosis: 'Price shock friction detected at checkout shipping field — driven by loss aversion.'

This is the step that no other analytics product takes. Any tool can produce a friction score. Only a system that connects ML explainability to a psychology framework can tell you what the score means and why it is happening.

Model 3 — Real-Time Conversion Predictor

The conversion predictor outputs a probability between 0.0 and 1.0 representing the likelihood that the current session will end in abandonment within the next two minutes. This score is updated every 30 seconds throughout the live session, allowing it to track the evolution of user intent over time.

For real-time updates, a lightweight logistic regression model — trained on the same feature set as the XGBoost models — runs alongside the full model to provide instantaneous updates as new events arrive. The XGBoost model runs at session end for the full, high-accuracy score. This two-tier approach balances real-time responsiveness with predictive accuracy.

Score range	Interpretation	Recommended merchant action
0.0 – 0.40	Low risk — user is engaged	No intervention needed, continue monitoring
0.40 – 0.65	Moderate friction detected	Surface social proof or reassurance elements
0.65 – 0.85	High abandonment risk	Trigger targeted intervention based on friction type
0.85 – 1.0	Critical — abandonment imminent	Immediate retention action — discount, guarantee, urgency

Layer 4 — The Suggestion Engine

The suggestion engine is the layer that transforms BehaviorIQ from a diagnostic tool into a strategic partner. Where the first three layers answer 'what is happening and why,' the suggestion engine answers 'what should the merchant do about it.'

Critically, the suggestion engine is not an ML model. It is a deterministic rules-based reasoning system built on top of the ML outputs. This is an intentional design decision: merchants will not act on recommendations they do not understand or trust, and explainability is therefore non-negotiable. Every recommendation the suggestion engine produces is traceable to a specific friction type, a specific psychological mechanism, a specific evidence base, and — where available — the merchant's own historical experiment data.

The engine operates in three steps. First, it receives the session intelligence object from the ML ensemble — friction type, psychological state, abandonment risk, affected element. Second, it queries the merchant's experiment history to determine whether similar friction has been tested before and what the results were. Third, it searches the intervention library — a curated database of psychology-backed fixes organized by friction type, with estimated conversion impact ranges derived from published research and aggregated experiment data — and returns a ranked list of recommendations.

> *Example output: 'Loss aversion detected on shipping cost field — hesitation score 3.4× above your baseline. Your experiment history contains no previous tests on this element. Research suggests reframing shipping cost as a savings threshold (e.g. "Add €12 more for free shipping") reduces perceived loss by anchoring attention on the gain. Predicted conversion improvement at this step: +6–11%, based on 23 comparable experiments in our database.'*

As the merchant's experiment history grows, the suggestion engine becomes progressively more personalized and accurate. A merchant with 50 logged experiments receives fundamentally better recommendations than one with 5 — because the engine can account for what has already been tried, what their specific audience responds to, and which psychological levers are most effective for their product category. This is the compounding moat that makes BehaviorIQ more valuable the longer a merchant stays.

Intelligence Layer — Technical Summary

Layer	Technology	Latency	Key output
Psychology framework	Structured JSON config, versioned	< 1ms lookup	Named psychological state
Feature engineering	Python worker, stateful session object	< 20ms per batch	Feature vector (8–12 floats)
ML ensemble (3 models)	XGBoost + SHAP, FastAPI endpoint	< 50ms per score	Intent, friction points, risk score
Suggestion engine	Rules-based, intervention library DB	< 30ms lookup	Ranked, justified test recommendations
End-to-end pipeline	Event batch → full intelligence output	< 300ms total	SessionIntelligence object

Market Opportunity

A large, growing, and underserved market

Market sizing

Market	Size (2026)	Growth rate	Our initial focus
Total addressable market (TAM)	€45B+	14% CAGR	Global e-commerce analytics & optimization
Serviceable addressable market (SAM)	€8.2B	18% CAGR	SMB to mid-market e-commerce (€100k–€50M revenue)
Serviceable obtainable market (SOM)	€180M	—	European e-commerce, Shopify/WooCommerce ecosystem

Why now

Three forces are converging to make this the right moment for BehaviorIQ. The first is the maturation of behavioral ML: transformer-based sequence models, gradient boosting on tabular data, and SHAP explainability have all reached a level of accessibility and performance that makes production-grade psychological inference feasible at startup scale. The second is the post-cookie world: as third-party cookies disappear and privacy regulations tighten, first-party behavioral data — collected with clear consent from sessions on owned properties — is becoming the most valuable and legally defensible form of customer intelligence. The third is the AI moment: merchants are primed to adopt AI-powered tools in a way they were not three years ago, lowering the educational barrier to our value proposition significantly.

Competitive landscape

Competitor	Strength	Weakness	Why we win
Hotjar	Brand recognition, ease of use	No ML, no psychology, no recommendations	Depth of insight is incomparably greater
Contentsquare	Enterprise relationships	Very expensive, complex, no psychology layer	Price + psychological intelligence advantage
FullStory	Strong session replay	No prediction, no suggestions, expensive	Real-time scoring + actionability
Optimizely	A/B testing depth	No behavioral intelligence, hypothesis-driven	Friction-first suggestions vs hypothesis guessing

AB Tasty / VWO	Mid-market A/B testing	No behavioral data, no psychology	Integrated behavioral intelligence layer

Business Model

How BehaviorIQ generates revenue

Pricing model — SaaS subscription by monthly sessions

BehaviorIQ is priced on a simple, transparent metric that scales directly with the value delivered: the number of monthly sessions tracked. This aligns our revenue growth directly with merchant growth and ensures that pricing feels proportional to benefit at every stage.

Tier	Monthly sessions	Price / month	Target customer	Key features
Starter	Up to 10,000	€49	Small merchants < €500k revenue	SDK, friction scores, dashboard, 30-day history
Growth	Up to 100,000	€149	Mid-size shops €500k–€5M revenue	All Starter + ML predictions, experiment memory, semantic search
Scale	Up to 500,000	€399	Larger merchants €5M–€30M revenue	All Growth + suggestion engine, API access, priority support
Enterprise	Unlimited	Custom	€30M+ revenue	All Scale + dedicated model, custom integrations, SLA

Revenue projections

Milestone	Timeframe	Merchants	Projected MRR	Projected ARR
First pilots	Month 1–3	10 (free/discounted)	€0 — validation phase	—
First revenue	Month 4–6	20 paying	€3,000–€5,000	€36k–€60k
Post-seed growth	Month 7–12	80 paying	€18,000–€25,000	€216k–€300k
Series A readiness	Month 18–24	300+ paying	€80,000–€120,000	€960k–€1.4M

Unit economics (at Growth tier, steady state)

Monthly recurring revenue per merchant	€149
Annual contract value (ACV)	€1,788
Estimated customer acquisition cost (CAC)	€200–€400 (content + outbound)

Estimated monthly churn (target)	< 3%
Payback period	2–3 months
Estimated LTV (at 3% monthly churn)	€4,967 (33-month avg lifetime)
LTV:CAC ratio	12:1 – 25:1 (target > 10:1)

Project Plan

From MVP to market — a phased execution roadmap

Phase 1 — MVP (Months 1–2): Prove the core thesis

The goal of Phase 1 is to build the minimum product that can prove a single hypothesis: will merchants install our SDK and find the friction data valuable enough to pay for? Everything that cannot directly prove or disprove this hypothesis is deferred.

Week	Deliverable	Technology	Success criterion
Week 1	Data preparation & ML baseline	Python, pandas, scikit-learn	Labeled dataset ready, XGBoost baseline trained
Week 2	ML model training & evaluation	XGBoost, SHAP, pytest	AUC > 0.72 on held-out test set
Week 3	JavaScript SDK v1	Vanilla JS, WebSocket	Captures 5 signals, < 15KB, async load
Week 4	Backend API + ML endpoint	FastAPI, Supabase, Railway	Events ingested, sessions scored, < 300ms
Week 5–6	Merchant dashboard	React, Recharts, Vercel	Friction heatmap, session list, risk scores visible
Week 7	Testing & hardening	pytest, Playwright, Sentry	All critical paths tested, error monitoring live
Week 8	First pilot merchant	Manual onboarding	Real sessions flowing, merchant confirms value

Phase 2 — Intelligence Deepening (Months 3–5): Build the moat

With pilot merchants generating real data, Phase 2 focuses on building the layers that make BehaviorIQ genuinely defensible: the full psychology framework, the SHAP-based friction detector, and the experiment memory system.

- Complete psychology framework: research and encode all 12 primary psychological states with behavioral signatures and intervention mappings
- Retrain ML models on real merchant data — replace bootstrap dataset with production sessions
- Build friction detector with SHAP explainability — produce human-readable friction diagnoses
- Launch experiment memory: merchants can log A/B test results with psychological tagging
- Implement semantic search on experiment history using sentence-transformers embeddings
- Build integrations: Google Sheets import, VWO and Optimizely result sync via API
- Onboard 10 pilot merchants — target mix of Shopify, WooCommerce, and custom shops

Phase 3 — Suggestion Engine & Growth (Months 6–9): The full product

Phase 3 delivers the suggestion engine — the layer that completes the loop from friction detection to intervention recommendation — and shifts focus toward scalable merchant acquisition.

- Build suggestion engine: intervention library, rules-based recommendation system, predicted impact scoring
- Launch self-serve onboarding: merchants can sign up, install SDK, and see data without manual support
- Shopify App Store submission — reach 4.8M+ Shopify merchants
- Content marketing engine: publish psychology-of-conversion research to drive organic discovery
- Target 80 paying merchants by end of Month 9
- Infrastructure scaling: move events table to ClickHouse for high-volume analytics

Phase 4 — Series A Preparation (Months 10–18): Scale

- Reach €80k MRR and 300+ paying merchants
- Publish first proprietary research report on e-commerce conversion psychology — drives press and enterprise leads
- Enterprise tier launch with dedicated model training per merchant
- Expand to additional e-commerce platforms: Magento, BigCommerce, headless commerce
- Hire ML engineer and second full-stack developer
- Raise Series A of €2–3M to accelerate European and US expansion

Key milestones summary

Milestone	Target date	Success metric
Working MVP with ML scoring	Month 2	AUC > 0.72, first sessions scored
10 pilot merchants onboarded	Month 3	Real data flowing, qualitative validation
First €5k MRR	Month 6	20 paying merchants at avg €250/mo
Suggestion engine live	Month 8	Merchants acting on recommendations
Shopify App Store listing	Month 7	Self-serve acquisition channel open
€25k MRR	Month 12	80+ paying merchants
Series A ready	Month 18	€80k MRR, 300+ merchants, strong NPS

Technology Stack

Built for speed, cost efficiency, and production reliability

MVP stack — optimized for near-zero cost

Layer	Technology	Rationale	Monthly cost
SDK	Vanilla JavaScript	No dependencies, < 15KB, works on any website	€0
Backend API	FastAPI (Python)	Fast, async, native Python ML ecosystem integration	—
Backend hosting	Railway	One-click Python deployment, free tier available	~€5
Database	Supabase (PostgreSQL)	Managed Postgres + auth + row-level security built in	€0 (free tier)
Frontend	React + Recharts	Component ecosystem, Recharts for data visualizations	—
Frontend hosting	Vercel	Global CDN, zero-config React deployment	€0 (free tier)
ML training	Python + XGBoost + SHAP	Industry standard, CPU-only training, < 5min per run	€0 (local)
ML serving	FastAPI endpoint on Railway	Model loaded at startup, < 10ms inference	Included above
Error monitoring	Sentry	Real-time error alerts, free tier covers MVP volume	€0
CI/CD	GitHub Actions	Automated tests and deployment on every push	€0
Semantic search	sentence-transformers	Open source, runs on CPU, no API cost	€0
Total running cost	—	—	~€5–10/month

Scale stack — after 100 merchants

Component	Upgrade	Trigger	Estimated cost
Events storage	Supabase → ClickHouse	Events table > 10M rows	€50–100/mo
Backend hosting	Railway → AWS ECS	API latency > 200ms p95	€100–300/mo
ML serving	Single endpoint → model registry	Multiple merchant-specific models	€50/mo

Search	Local embeddings → Pinecone	Experiment DB > 10,000 entries	€70/mo

Use of Funds

How the €350,000 pre-seed round will be deployed

Allocation

Category	Amount	% of raise	What it buys
Product development	€175,000	50%	Technical co-founder runway (18 months), infrastructure, tooling, security audit
Team	€105,000	30%	ML engineer hire (Month 4), part-time designer (Month 3), part-time growth (Month 6)
Go-to-market	€70,000	20%	Content marketing, Shopify App Store launch, conference presence, outbound sales tooling

The €350,000 raise is designed to take BehaviorIQ from working MVP to €25k MRR with 80+ paying merchants — the metrics required to raise a Series A of €2–3M at a significantly higher valuation.

18-month runway plan

Month 1–3	Technical co-founder builds MVP solo with AI-assisted development. Zero external hires. Infrastructure cost < €50/mo.
Month 4	Hire first ML engineer (part-time, €3,500/mo) to accelerate intelligence layer development.
Month 5–6	First paying merchants. Begin content marketing. Hire part-time designer.
Month 9	Shopify App Store launch. Begin outbound sales motion.
Month 12	80 paying merchants. €25k MRR. Begin Series A process.
Month 18	Series A close. 300+ merchants. €80k MRR.

Vision & Long-Term Ambition

What BehaviorIQ becomes in five years

The 5-year vision

BehaviorIQ's immediate product — real-time behavioral intelligence for e-commerce conversion optimization — is the first expression of a much larger platform vision. The fundamental insight that human behavior on digital surfaces is rich with psychological signal, and that AI can now decode those signals in real time, applies far beyond checkout optimization.

In five years, BehaviorIQ is the behavioral intelligence infrastructure layer for digital commerce globally — the system that every e-commerce operator runs because understanding the psychology of their customers has become as fundamental as knowing their revenue. The platform has expanded from friction detection to a complete psychological customer intelligence system: knowing not just what a user did on a specific session, but building a longitudinal model of their decision-making psychology over time.

Expansion roadmap

Horizon	Timeframe	Product expansion
H1 — E-commerce intelligence	Now → Year 2	Friction detection, experiment memory, suggestion engine for Shopify/WooCommerce
H2 — Cross-platform behavioral data	Year 2–3	Behavioral intelligence for mobile apps, email, ads — any digital touchpoint
H3 — Predictive customer psychology	Year 3–4	Longitudinal psychological profiles per customer segment — not just session-level signals
H4 — Behavioral intelligence API	Year 4–5	Platform opens as API — third parties build on BehaviorIQ psychological intelligence

The defensibility story

BehaviorIQ becomes more defensible with every passing month through three compounding moats. The first is data: as more merchants use the platform, the training data for our models grows richer, making our predictions more accurate than any new entrant could achieve. The second is experiment memory: each merchant's knowledge base of what works for their specific audience is stored exclusively in BehaviorIQ — switching costs grow with every experiment logged. The third is the psychology framework: as our research team extends the framework with findings from our own data, it becomes a proprietary body of knowledge that is increasingly difficult and expensive to replicate.

The founding team's unfair advantage

The technical co-founder brings a rare combination of full-stack engineering capability, machine learning competency, and an understanding of the behavioral economics literature that directly underpins the product. In a market full of analytics tools built by engineers who don't understand psychology, and psychology tools built by researchers who don't understand engineering, BehaviorIQ is built by someone who can hold both simultaneously.

The business co-founder brings deep domain expertise in e-commerce and behavioral intelligence, the Frankfurt network to access early pilot merchants, and the go-to-market experience to translate a technically complex product into a commercially compelling proposition.

> *We are not building a better heatmap tool. We are building the first platform that gives e-commerce operators a genuine understanding of their customers' psychology — and the systematic capability to act on that understanding. The market is large, the timing is right, the technology is ready, and the team is uniquely positioned to execute. We invite you to build it with us.*

Contact

Company	BehaviorIQ (pre-incorporation)
Stage	Pre-seed
Location	Frankfurt am Main, Germany
Document version	1.0 — March 2026
Status	Confidential — do not distribute